

Analyze A/B Test Results ¶

Table of Contents

- [Introduction](#)
- [Part I - Probability](#)
- [Part II - A/B Test](#)
- [Part III - Regression](#)

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

Part I - Probability

To get started, let's import our libraries.

In [1]:

```
import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
import warnings
warnings.filterwarnings('ignore')
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. **Use your dataframe to answer the questions in Quiz 1 of the classroom.**

a. Read in the dataset and take a look at the top few rows here:

In [2]:

```
df=pd.read_csv("ab_data.csv")
df.head(2)
```

Out[2]:

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0

b. Use the below cell to find the number of rows in the dataset.

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id         294478 non-null  int64
1   timestamp       294478 non-null  object
2   group           294478 non-null  object
3   landing_page    294478 non-null  object
4   converted       294478 non-null  int64
dtypes: int64(2), object(3)
memory usage: 7.9+ MB
```

c. The number of unique users in the dataset.

In [4]:

```
df.user_id.nunique()
```

Out[4]:

```
290584
```

d. The proportion of users converted.

In [5]:

```
df.converted.mean()
```

Out[5]:

```
0.11965919355605512
```

e. The number of times the `new_page` and `treatment` don't line up.

In [6]:

```
df.query('group == "treatment" and landing_page == "old_page" or group == "control" and lan
```

Out[6]:

```
3893
```

f. Do any of the rows have missing values?

In [7]:

```
df.isnull().sum()
#or df.info()
```

Out[7]:

```
user_id      0
timestamp    0
group        0
landing_page  0
converted    0
dtype: int64
```

2. For the rows where **treatment** is not aligned with **new_page** or **control** is not aligned with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to provide how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

In [8]:

```
df2 = df[((df.group == 'control') & (df.landing_page == 'old_page')) | ((df.group == 'treat
```

In [9]:

```
df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].shape[
```

Out[9]:

0

In [10]:

```
print(df2.shape)
df2.head()
```

(290585, 5)

Out[10]:

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

3. Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_ids** are in **df2**?

In [11]:

```
df2.user_id.nunique()
```

Out[11]:

290584

b. There is one **user_id** repeated in **df2**. What is it?

In [12]:

```
df2[df2.user_id.duplicated()]
```

Out[12]:

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

c. What is the row information for the repeat **user_id**?

In [13]:

```
df2[df2.user_id==773192]
```

Out[13]:

	user_id	timestamp	group	landing_page	converted
1899	773192	2017-01-09 05:37:58.781806	treatment	new_page	0
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

In [14]:

```
df2.drop(1899,inplace=True)
```

4. Use **df2** in the below cells to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

In [15]:

```
df2['converted'].mean()
```

Out[15]:

```
0.11959708724499628
```

b. Given that an individual was in the `control` group, what is the probability they converted?

In [16]:

```
df2.query('group == "control"')['converted'].mean()
```

Out[16]:

```
0.1203863045004612
```

c. Given that an individual was in the `treatment` group, what is the probability they converted?

In [17]:

```
df2.query('group == "treatment"')['converted'].mean()
```

Out[17]:

```
0.11880806551510564
```

d. What is the probability that an individual received the new page?

In [18]:

```
(df2[df2.landing_page=="new_page"].user_id.count())/df2.shape[0]
```

Out[18]:

```
0.5000619442226688
```

e. Consider your results from a. through d. above, and explain below whether you think there is sufficient evidence to say that the new treatment page leads to more conversions.

The probability of converting regardless of page is: 11.96%. The probability of converting from control group is 12.03%. The probability of an individual converting from treatment group is 11.88%. The probability of receiving the new page is: 50.01%

The probability of converting in the control group and the treatment group are very close, with a difference of only 0.16%. So we need to more evidence to prove.

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new} , which are the converted rates for the old and new pages.

$H_0: p(\text{old}) - p(\text{new}) \geq 0$ $H_1: p(\text{old}) - p(\text{new}) < 0$

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the **convert rate** for p_{new} under the null?

In [19]:

```
p_new = df2.converted.mean()  
p_new
```

Out[19]:

0.11959708724499628

b. What is the **convert rate** for p_{old} under the null?

In [20]:

```
p_old = df2.converted.mean()  
p_old
```

Out[20]:

0.11959708724499628

c. What is n_{new} ?

In [21]:

```
n_new = df2.query('landing_page == "new_page"')['user_id'].count()
print(n_new)
```

145310

d. What is n_{old} ?

In [22]:

```
n_old = df2.query('landing_page == "old_page"')['user_id'].count()
print(n_old)
```

145274

e. Simulate n_{new} transactions with a convert rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

In [23]:

```
new_page_converted = np.random.choice([0,1],n_new, p=(p_new,1-p_new))
new_page_converted
```

Out[23]:

array([1, 1, 1, ..., 1, 1, 1])

f. Simulate n_{old} transactions with a convert rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

In [24]:

```
old_page_converted = np.random.choice([0,1],n_new, p=(p_old,1-p_old))
old_page_converted
```

Out[24]:

array([1, 1, 1, ..., 1, 1, 1])

g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).

In [25]:

```
new_mean = new_page_converted.mean()
old_mean = old_page_converted.mean()
print(new_mean - old_mean)
```

0.0014933590255316043

h. Simulate 10,000 $p_{new} - p_{old}$ values using this same process similarly to the one you calculated in parts a.

through **g.** above. Store all 10,000 values in a numpy array called **p_diffs**.

In [26]:

```
p_diffs = []
for _ in range(10000):
    new_page_converted = np.random.choice([0,1], size=n_new, p=[1-p_new, p_new])
    old_page_converted = np.random.choice([0,1], size=n_old, p=[1-p_old, p_old])
    p_diffs.append(new_page_converted.mean() - old_page_converted.mean())
```

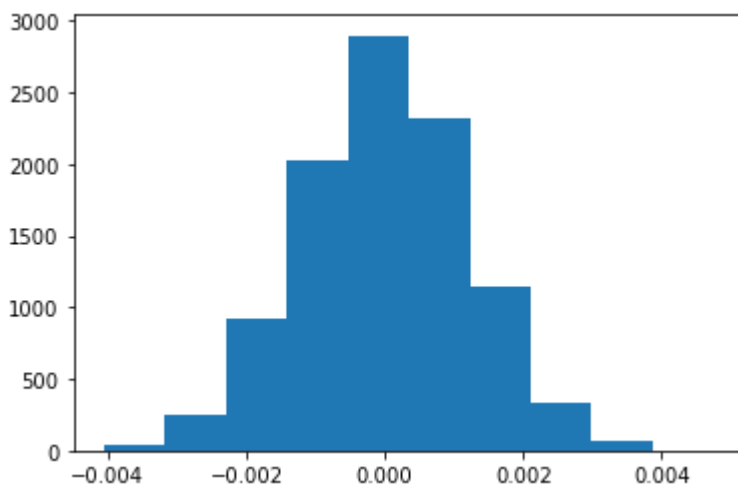
i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

In [27]:

```
plt.hist(p_diffs)
```

Out[27]:

```
(array([ 37., 250., 920., 2031., 2897., 2324., 1145., 330., 62.,
        4.]),
 array([-0.00405587, -0.00317559, -0.00229531, -0.00141502, -0.00053474,
        0.00034554, 0.00122583, 0.00210611, 0.00298639, 0.00386668,
        0.00474696]),
 <a list of 10 Patch objects>)
```



j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

In [28]:

```
convert_new = df2[(df2.converted==1) & (df.landing_page == "new_page")].user_id.nunique()
convert_old = df2[(df2.converted==1) & (df.landing_page == "old_page")].user_id.nunique()
actual_new=convert_new/ n_new
actual_old=convert_old/n_old
obs_diff = actual_new - actual_old
obs_diff
```

Out[28]:

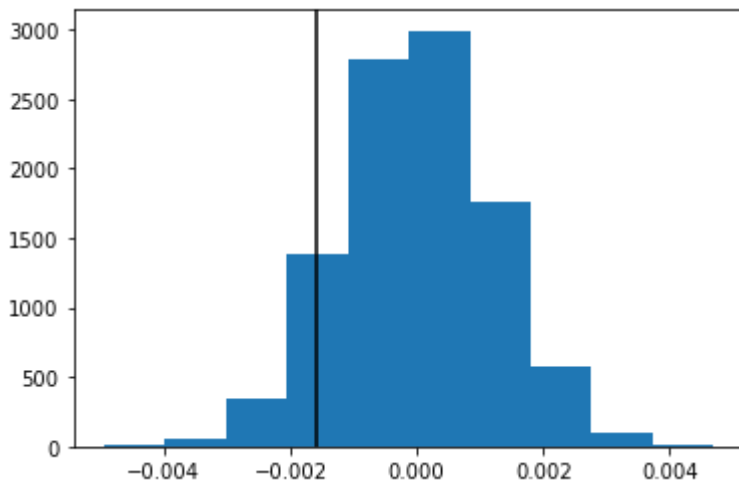
```
-0.0015782389853555567
```


In [29]:

```

null_vals = np.random.normal(0, np.std(p_diffs), np.array(p_diffs).size)
plt.hist(null_vals)
plt.axvline(x=obs_diff,color='black');

```



In [30]:

```

(null_vals > obs_diff).mean()

```

Out[30]:

0.9062

k. In words, explain what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

On calculation of P value came as 0.904 i.e more than 90.4% of our sample values is above the observed difference. From our histogram and results, the new page does not do better and we fail to reject the null hypothesis. Concluding: With type I error (rate = 0.05) the old page has higher probability of converting.

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer the the number of rows associated with the old page and new pages, respectively.

In [31]:

```

import statsmodels.api as sm

convert_old = df2.query('landing_page == "old_page" and converted == 1')['user_id'].count()
convert_new = df2.query('landing_page == "new_page" and converted == 1')['user_id'].count()
n_old = df2.query('landing_page == "old_page"')['user_id'].count()
n_new = df2.query('landing_page == "new_page"')['user_id'].count()

```

m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. [Here \(http://knowledgetack.com/python/statsmodels/proportions_ztest/\)](http://knowledgetack.com/python/statsmodels/proportions_ztest/) is a helpful link on using the built in.

In [32]:



```
z_score, p_value = sm.stats.proportions_ztest([convert_new, convert_old], [n_new, n_old], a
print('Z:', z_score)
print('P:', p_value)
```

Z: -1.3109241984234394

P: 0.9050583127590245

In [33]:



```
from scipy.stats import norm
z_critical=norm.ppf(1-(0.05))
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts **j.** and **k.**?

The p value obtained here is similar to that obtained in j part. The z score is also less than critical z score. We can now accept the null hypothesis that the conversion rates of the old page are equal or better than the conversion rates of the new page.

Part III - A regression approach

1. In this final part, you will see that the result you achieved in the previous A/B test can also be achieved by performing regression.

a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

As each row can have only two possibilities: a conversion or no conversion So we can use a logistic regression here.

b. The goal is to use **statsmodels** to fit the regression model you specified in part **a.** to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

In [34]:



```
df2.head(2)
```

Out[34]:

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0

In [35]:

```
df2['intercept'] = 1
```

In [36]:

```
df2[['control', 'ab_page']] = pd.get_dummies(df2['group'])  
df2.drop(['control'], axis=1, inplace=True)  
df2.head(2)
```

Out[36]:

	user_id	timestamp	group	landing_page	converted	intercept	ab_page
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	1	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	1	0

c. Use **statsmodels** to import your regression model. Instantiate the model, and fit the model using the two columns you created in part **b.** to predict whether or not an individual converts.

In [37]:

```
logit_mod = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])  
res = logit_mod.fit()
```

```
Optimization terminated successfully.  
Current function value: 0.366118  
Iterations 6
```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

In [38]:



```
res.summary()
```

Out[38]:

Logit Regression Results

Dep. Variable:	converted	No. Observations:	290584
Model:	Logit	Df Residuals:	290582
Method:	MLE	Df Model:	1
Date:	Sun, 31 May 2020	Pseudo R-squ.:	8.077e-06
Time:	22:07:48	Log-Likelihood:	-1.0639e+05
converged:	True	LL-Null:	-1.0639e+05
Covariance Type:	nonrobust	LLR p-value:	0.1899

	coef	std err	z	P> z	[0.025	0.975]
intercept	-1.9888	0.008	-246.669	0.000	-2.005	-1.973
ab_page	-0.0150	0.011	-1.311	0.190	-0.037	0.007

e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**?

Hint: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in the **Part II**?

The p value associated with **ab_page** is 0.19. The value differs from the one found in part-2 as the hypothesis were different. For the second part we calculate the probability receiving a observed statistic if the null hypothesis is true. Therefore this is a one-sided test, whereas the **ab_page** p-value is a result of a two sided test. The null hypothesis for this case is : there should be no significant relationship between the conversion rate and page.

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

The regression model could have other influential factors like : Educational background, time spend browsing the website. Even if he has not converted but spend a reasonable amount of time on the website , he might convert later or suggest others (It could be a positive sign).

However many additional terms could complex the process and thereby increasing the difficulty on analysing final results.

g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives. You will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows. [Here \(https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.join.html\)](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.join.html) are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables.** Provide the statistical output as well as a written response to answer this question.

In [39]:

```
countries_df = pd.read_csv('./countries.csv')
df_new = countries_df.set_index('user_id').join(df2.set_index('user_id'), how='inner')
```

In [40]:

```
countries_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 290584 entries, 0 to 290583
Data columns (total 2 columns):
#   Column   Non-Null Count  Dtype
---  -
0   user_id  290584 non-null  int64
1   country  290584 non-null  object
dtypes: int64(1), object(1)
memory usage: 3.3+ MB
```

In [41]:

```
countries_df.nunique()
```

Out[41]:

```
user_id    290584
country         3
dtype: int64
```

In [42]:

```
countries_df.country.unique()
```

Out[42]:

```
array(['UK', 'US', 'CA'], dtype=object)
```

In [43]:

```
df_new[['CA', 'UK', 'US']] = pd.get_dummies(df_new['country'])
```

In [44]:

```
df_new = df_new.drop(['CA'], axis=1)
```

h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

In [45]:



```
df_new.head()
```

Out[45]:

	country	timestamp	group	landing_page	converted	intercept	ab_page	UK
user_id								
834778	UK	2017-01-14 23:08:43.304998	control	old_page	0	1	0	1
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	0	1	1	0
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	1	1	1	1
711597	UK	2017-01-22 03:14:24.763511	control	old_page	0	1	0	1
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	0	1	1	1

In [46]:



```
logit_mod = sm.Logit(df_new['converted'], df_new[['intercept', 'ab_page', 'US', 'UK']])  
results = logit_mod.fit()
```

```
Optimization terminated successfully.  
Current function value: 0.366113  
Iterations 6
```

In [47]:



```
results.summary()
```

Out[47]:

Logit Regression Results

Dep. Variable:	converted	No. Observations:	290584
Model:	Logit	Df Residuals:	290580
Method:	MLE	Df Model:	3
Date:	Sun, 31 May 2020	Pseudo R-squ.:	2.323e-05
Time:	22:07:53	Log-Likelihood:	-1.0639e+05
converged:	True	LL-Null:	-1.0639e+05
Covariance Type:	nonrobust	LLR p-value:	0.1760

	coef	std err	z	P> z	[0.025	0.975]
intercept	-2.0300	0.027	-76.249	0.000	-2.082	-1.978
ab_page	-0.0149	0.011	-1.307	0.191	-0.037	0.007
US	0.0408	0.027	1.516	0.130	-0.012	0.093
UK	0.0506	0.028	1.784	0.074	-0.005	0.106

As the P-values are larger than 0.005 so we fail to reject null hypothesis. Therefore countries don't effect the conversion rate.

Conclusions

New Page could not result in more conversion rate,so we should keep the old page.

After computing by different techniques the old model looks more profitable.

The countries did not play a significant role in terms of conversion rate.